File System Interpretation

INTRODUCTION:

We are all familiar with the characteristics of the files and directories in which we store all of our data. As with many other persistent objects, their functionality, generality, performance and robustness all derive from the underlying data structures in terms of which they are implemented. In this project we will design and implement a program to read the on-disk representation of a file system, analyze it, and summarize its contents (in several csv files). In the next project, we will write a program to analyze this summary for evidence of corruption.

This project can be broken into two major steps:

- Understand the on-disk data format of the EXT2 file system. We will mount a provided image file on your own Linux and explore it with familiar file navigation commands and *debugfs(8)*.
- Write a program to analyze the file system in that image file and output a summary to six csv files (describing the super block, groups, free-lists, inodes, indirect blocks, and directories). This may involve much more code than we have written in previous projects, but (once the underlying data structures are understood) the actual code will be fairly simple.

RELATION TO READING AND LECTURES:

This project more deeply explores the filesystem structures described in Arpaci chapter 39.
The images we will be working with are EXT2 file systems, as described in sections 40.2-40.5.

PROJECT OBJECTIVES:

- Primary: reinforce the basic file system concepts of directory objects, file objects, and free space.
- Primary: reinforce the implemenation descriptions provided in the text and lectures.
- Primary: gain experience examining, interpreting and processing information in complex binary data structures.
- Secondary: gain practical experience with typical on-disk file system data formats.

DELIVERABLES:

A single tarball (`.tar.gz`) containing:

- (at least) one C/C++ source module that compiles cleanly with no errors or warnings).

- A `Makefile` to build and run the deliverable program. The higher level targets should be:
    - default ... compile your program to produce an executable named `lab3a`
    - **dist** ... create the deliverable tarball
    - **clean** ... delete all generated programs and output, returning directory to its freshly un-tar-ed state.

- a `README` text file containing descriptions of each of the included files and any other information about your submission that you would like to bring to our attention (e.g. research, limitations, features, testing methodology, use of slip days).

PROJECT DESCRIPTION:

Historically, file systems were almost always been implemented as part of the operating system, running in kernel mode. Kernel code is expensive to develop, difficult to test, and prone to catastrophic failures. Within the past 15 years or so, new developments have made it possible to implement file systems in user mode, improving maintainability ... and in some cases delivering even better performance than could be achieved with kernel code. All of this project will be done as user-mode software.

To ensure data privacy and integrity, file system disks are generally protected from access by ordinary applications. Linux supports the creation, mounting, checking, and debugging of file systems stored in ordinary files. In this project, we will provide EXT2 file system images in ordinary files. Because they are in ordinary files (rather than protected disks) you can access/operate on those file system images with ordinary user mode code.

PREPARATION

To perform this assignment, you may need to study a few things:

- debugfs(8). a tool for exploring on-disk file system structures.
- pread(2), an alternative to *read(2)* for random-access file processing.
- a comprehensive overview of the EXT2 file system format.
- a slightly simplified version of the Linux header file that defines the format of the EXT2 file system. Please use this header file and include it in your submission. Do not assume that the standard header file will be available on the test system.

PART 1: Exploring an EXT2 Image

In order to mount and explore a file system (even one stored in an ordinary file) you will need the ability to run privileged commands (e.g. *mount(8)*). Since you will not have *sudo(8)* access on departmental servers, you will have to do this exploration on your own personal Linux system.

Download this (~2MB) file system image, and mount it onto your own Linux, with the following commands:

```
mkdir fs
sudo mount –o loop ext2_image fs
sudo chown –R USER fs          # where USER is your user-name
```

Note: this file is 200 Mbytes in size.

Now, you can navigate the file system, just like an ordinary directory, with commands like *ls(1)*, *cat(1)*, and *cd(1)*. After you are done with it, you can unmount with the following command:

```
sudo umount fs
```

Before you start writing your C/C++ program to interpret the diskimage file, you can explore it further using *debugfs(8)* (on your own Linux system). You will find many useful commands in its man page. Some particularly helpful ones are: `stats`, `stat`, `bd`, `testi`, and `testb`. If you have problems interpreting the file system (or being sure that you have correctly interpreted it), you can use the *debugfs* program to help you understand the correct contents.

To ensure you are correctly interpreting the file system image, we have included many unusual things, which might not be noticed by a naive implementation:

- sparse files
- very large files
- allocated data blocks full of zeroes
- unallocated blocks containing valid data

- files with data beyond their length
- files with long names
- files with syntactically strange names
- directories that span multiple blocks and have obsolete entries for deleted files

PART 2: Summarizing an EXT2 Image

In this step, you will write a C/C++ program called `lab3a` that:

- Reads a file system image, whose name is specified as a command line argument. For example, we may run your program with the above file system image using the a command like:

    ```
    ./lab3a EXT2_test.img
    ```

- Analyzes the provided file system image and produces (to standard out) CSV summaries of what it finds. The contents of these CSV lines described below. Your program must output these files with exactly the same formats as shown below. We will use *sort(1)* and *diff(1)* to compare your csv output with ours, so a different format will make your program fail the test.

Please note that, although you cannot mount the provided image file and run *debugfs* on departmental servers, your lab3a program should, like previous assignments, be able to run on departmental servers.

There are six types of output lines that your program should produce, each summarizing a different part of the file system. Remember, you can always check your program's output against *debugfs*'s output. All the information required for the summary can be manually found and checked by using debugfs. We have also included (for testing purposes) a much smaller image as well as a correct summary.

You are free to produce additional commentary to stderr, but only file system summary information should be logged to stdout.

superblock summary

A single new-line terminated line, comprised of eight comma-separated fields (with no white-space), summarizing the key file system parameters:

1. SUPERBLOCK
2. total number of blocks (decimal)
3. total number of i-nodes (decimal)
4. block size (in bytes, decimal)
5. i-node size (in bytes, decimal)
6. blocks per group (decimal)
7. i-nodes per group (decimal)
8. first non-reserved i-node (decimal)

group summary

Scan each of the groups in the file system. For each group, produce a new-line terminated line for each group, each comprised of nine comma-separated fields (with no white space), summarizing its contents.

1. GROUP
2. group number (decimal, starting from zero)
3. total number of blocks in this group (decimal)
4. total number of i-nodes in this group (decimal)

5. number of free blocks (decimal)
6. number of free i-nodes (decimal)
7. block number of free block bitmap for this group (decimal)
8. block number of free i-node bitmap for this group (decimal)
9. block number of first block of i-nodes in this group (decimal)

Note that most Berkeley-derived file systms (like EXT2) support blocks, and fragments, which may have different sizes. The block is the preferred unit of allocation. But in some cases, fragments may be used (to reduce internal fragmentation loss). Block addresses and the free block list entries are based on the fragment size, rather than the block size. But, in the images we give you, the block and fragment sizes will be the same.

It is also the case that most Berkeley-derived file systems put a copy of the superblock and group summary at the start of each group. But, in the images we give you, there will only be one group.

free block entries

Scan the free block bitmap for each group. For each free block, produce a new-line terminated line, with two comma-separated fields (with no white space).

1. BFREE
2. number of the free block (decimal)

Take care to verify that you:

1. understand whether 1 means allocated or free.
2. have correctly understood the block number to which the first bit corresponds.
3. do not interpret more bits than there are blocks in the group.

free I-node entries

Scan the free I-node bitmap for each group. For each free I-node, produce a new-line terminated line, with two comma-separated fields (with no white space).

1. IFREE
2. number of the free I-node (decimal)

Take care to verify that you:

1. understand whether 1 means allocated or free.
2. have correctly understood the I-node number to which the first bit corresponds.
3. do not interpret more bits than there are I-nodes in the group.

I-node summary

Scan the I-nodes for each group. For each valid (non-zero mode and non-zero link count) I-node, produce a new-line terminated line, with 27 comma-separated fields (with no white space). The first twelve fields are i-node attributes:

1. INODE
2. inode number (decimal)
3. file type ('f' for file, 'd' for directory, 's' for symbolic link, '?" for anything else)
4. mode (octal)
5. owner (decimal)
6. group (decimal)

7. link count (decimal)
8. creation time (mm/dd/yy hh:mm:ss, GMT)
9. modification time (mm/dd/yy hh:mm:ss, GMT)
10. time of last access (mm/dd/yy hh:mm:ss, GMT)
11. file size (decimal)
12. number of blocks (decimal)

The next fifteen fields are block addresses (decimal, 12 direct, one indirect, one double indirect, one tripple indirect).

directory entries

For each directory I-node, scan every data block. For each valid (non-zero I-node number) directory entry, produce a new-line terminated line, with seven comma-separated fields (no white space).

1. `DIRENT`
2. parent inode number (decimal) ... the I-node number of the directory that contains this entry
3. byte offset (decimal) of this entry within the directory
4. inode number of the referenced file (decimal)
5. entry length (decimal)
6. name length (decimal)
7. name (string, surrounded by single-quotes). Don't worry about escaping, we promise there will be no single-quotes or commas in any of the file names.

indirect block references

The I-node summary contains a list of all 12 blocks, and the primary single, double, and tripple indirect blocks. We also need to know about the blocks that are pointed to by those indirect blocks. For each file or directory I-node, scan the single indirect blocks and (recursively) the double and tripple indirect blocks. For each non-zero block pointer you find, produce a new-line terminated line with six comma-separated fields (no white space).

1. `INDIRECT`
2. I-node number of the owning file (decimal)
3. (decimal) level of indirection for the block being scanned ... 1 single indirect, 2 double indirect, 3 tripple
4. file offset (decimal) represented by the referenced block. If the referenced block is a data block, this is the logical block offset of that block within the file. If the referenced block is a single- or double-indirect block, this is the same as the logical offset of the first data block to which it refers.
5. block number of the (1,2,3) indirect block being scanned (decimal) ... not the highest level block (in the recursive scan), but the lower level block that contains the block reference reported by this entry.
6. block number of the referenced block (decimal)

If an I-node contains a tripple indirect block:

- the tripple indirect block number would be included in the INODE summary.
- INDIRECT entries (with level 3) would be produced for each double indirect block pointed to by that tripple indirect block.
- INDIRECT entries (with level 2) would be produced for each indirect block pointed to by one of those double indirect blocks.
- INDIRECT entries (with level 1) would be produced for each data block pointed to by one of those indirect blocks.

Sample Output

We have provided a very simple test file system image (that you can download and test with) along with correct summary output. Your program should be able to generate (modulo line ordering) the same output.

SUMMARY OF EXIT CODES:

- 0 ... analysis successful
- 1 ... bad arguments
- 2 ... corruption detected or other processing errors

SUBMISSION:

Your **README** file must include lines of the form:

NAME: *your name*
EMAIL: *your email*
ID: *your student ID*

Your name, student ID, and email address should also appear as comments at the top of your `Makefile` and each source file.

Your tarball should have a name of the form `lab3a-`*studentID*`.tar.gz`.
You can sanity check your submission with this test script.
Projects that do not pass the test script will not be accepted!

We will test it on a departmental Linux server. You would be well advised to test your submission on that platform before submitting it.

GRADING:

Points for this project will be awarded:

value	feature
	Packaging and build (10% total)
3%	un-tars expected contents
3%	clean build with default action (no warnings)
2%	correct `clean` and `dist` targets
2%	reasonableness of `README` contents
	Results (85% total)
10%	superblock summary
10%	group summaries
10%	free block entries
10%	free I-node entries
15%	I-node summaries
15%	directory entries
15%	block references
	Code Review (5%)
5%	general organization and readability